

# ROB311 Quiz 3

Hanhee Lee

April 3, 2025

## Contents

<b>1</b>	<b>Monte-Carlo Tree Search (MCTS)</b>	<b>2</b>
1.1	1 Player vs. 2 Player Turn-Taking Game Tree . . . . .	4
1.2	Examples . . . . .	5
1.2.1	1 Player Turn-Taking Game Tree . . . . .	5
1.2.2	2 Player Turn-Taking Game Tree . . . . .	6

# Turn-Taking Multi-Agent Decision Algorithms

## 1 Monte-Carlo Tree Search (MCTS)

### Algorithm:

1. Selection: Traverse using an alternate policy until a node has unexplored children.




Figure 1

- Our Agent (Upper Triangle): Uses UCB to choose the next node to explore
  - Other Agent (Down Triangle): Can't control their actions, so this agent picks w/ their own heuristic.
  - Square Boxes: Estimated values (i.e.  $n$  and  $\hat{q}$ )
  - Ends when there is at least one action that hasn't been explored yet. In this case, two actions haven't been explored.
  - Can skip expansion and simulation if the most recently expanded node is a terminal state.
2. Expansion: Expand an unexplored child; initialize  $n(a)$  and  $\hat{q}(s, a)$ .




Figure 2

- $\hat{q}(s, a)$  is initialized to 0 and  $n(a)$  is initialized to 1 b/c we've visited this node once.
  - Randomly pick an unexplored action unless there is only one action left.
  - Can skip simulation if the most recently expanded node is a terminal state.
3. Simulation: Traverse using the random policy until a terminal node is reached.




Figure 3

- Using random policy to simulate the game until a terminal state is reached (i.e. reward is obtained)
4. Back-propagation: Get the reward and reverse; update  $n(a)$  and  $\hat{q}(s, a)$ .




Figure 4

- Go up the path in yellow and update the values of  $n(a)$  and  $\hat{q}(s, a)$  for OUR agent only (i.e. the upper triangle)

## Warning:

- Works for more than 2 agents.
- Don't need to know anyone else's reward function.
- Has to be turn taking but can be not alternating (i.e. immediate switch between agents)
- Can augment simultaneous actions
- Communication
- Works for non-zero sum games.

## 1.1 1 Player vs. 2 Player Turn-Taking Game Tree

### Notes:

- 1 Player:  
—
- 2 Player:  
—

## 1.2 Examples

### 1.2.1 1 Player Turn-Taking Game Tree

### 1.2.2 2 Player Turn-Taking Game Tree

**Example:**

1. **Given:** Consider a simplified two-player turn-based game tree. You are currently at the root node  $S_0$ , which has three possible actions  $a_1, a_2, a_3$ . The current statistics of its children are as follows:

Action	$N(s_0, a)$	$\bar{X}(s_0, a)$
$a_1$	10	0.6
$a_2$	5	0.8
$a_3$	0	–

- $N(s_0, a)$ : Number of times action  $a$  has been selected at state  $s_0$
  - $\bar{X}(s_0, a)$ : Average reward obtained from action  $a$  at state  $s_0$
  - $UCB = \bar{X}(s_0, a) + \sqrt{\frac{\ln(t)}{N(s_0, a)}}$ 
    - $t$ : Total number of actions taken at  $s_0$
2. **Problems:**
    - If we were to use the UCB algorithm, which nodes get selected during the selection phase? Which node gets expanded during the expansion phase?
    - Suppose from the expanded node, simulation is performed until termination. A reward of +1 is obtained. Update the statistics at  $s_0$  accordingly.
    - Then, repeat the question, assuming a reward of –1 is attained after the simulation phase.
  3. **Solution:**
    - (a) **Selection 1:**  $s_0$  since we traverse until a node has unexplored children (i.e.  $s_3$  is unexplored)
    - (b) **Expansion 1:**  $s_3$  is automatically expanded since it is the only unexplored child of  $s_0$  w/  $N(s_0, a_3) = 1$  and  $\bar{X}(s_0, a_3) = 0$
    - (c) **Simulation 1:** Get a reward of +1
    - (d) **Back Propagation 1:** For this edge from  $s_0$  to  $s_3$ , we update the statistics as follows:
      - $N(s_0, a_3) = 1$
      - $\bar{X}(s_0, a_3) = \frac{1}{1} = 1$
    - (e) **Selection 2:**  $s_0$  and choose the action with the highest UCB value for  $s_1, s_2$ , and  $s_3$ :
      - $UCB(s_0, a_1) = 0.6 + \sqrt{\frac{\ln(16)}{10}} = 1.13$
      - $UCB(s_0, a_2) = 0.8 + \sqrt{\frac{\ln(16)}{5}} = 1.54$
      - $UCB(s_0, a_3) = 1 + \sqrt{\frac{\ln(16)}{1}} = 2.67$ . Therefore, choose  $s_3$  as part of the selection phase and assume it has unexplored children.
    - (f) **Expansion 2:** Not enough info but assume we expand an unexplored child.
    - (g) **Simulation 2:** Get a reward of –1
    - (h) **Back Propagation 2:** For this edge from  $s_0$  to  $s_3$ , we update the statistics as follows:
      - $N(s_0, a_3) = 2$
      - $\bar{X}(s_0, a_3) = \frac{1 + (-1)}{2} = 0$

**Example:**

1. **Given:** Consider (partial) 2-player turn-taking game-tree in which 21 iterations of MCTS have already been performed:

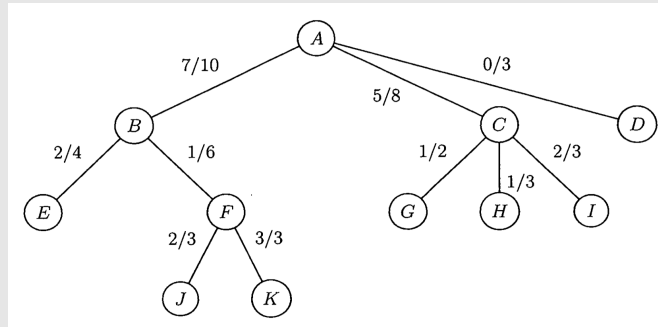


Figure 5

- Total reward: Numerator
  - Total number of times action  $a$  has been selected at state  $s$ : Denominator
2. **Problem:** If we use UCB to rank order state-action pairs, which of the following states will be chosen during the 22nd selection phase.

3. **Solution:**

- $UCB(AB) = 7/10 + \sqrt{\frac{\ln(21)}{10}} = 1.25$ 
  - $UCB(BE) = 2/4 + \sqrt{\frac{\ln(10)}{4}} = 1.26$
  - $UCB(BF) = 1/6 + \sqrt{\frac{\ln(10)}{6}} = 0.79$
- $UCB(AC) = 5/8 + \sqrt{\frac{\ln(21)}{8}} = 1.24$
- $UCB(AD) = 0/3 + \sqrt{\frac{\ln(21)}{3}} = 1.01$
- Therefore, choose A,B,E by selecting the nodes with the highest UCB values.

**Example:**

1. **Given:** Consider (partial) 2-player turn-taking game-tree in which 9 iterations of MCTS have already been performed:




Figure 6

- **Fix:**  $CG$  has 0/2 not 0/1 and  $CH$  has 0/1 not 1/1
2. **Problem:** Suppose path chosen during the 10th selection phase had the state sequence  $\langle A, C, H \rangle$  (i.e.  $H$  is the state expanded during the 10th expansion phase)
    - The simulation phase lasts for 12 transitions, after which a terminal state is reached.
    - The reward to the last turn-taker was +4.
    - Find  $q(A, \langle A, B \rangle)$ ,  $q(A, \langle A, C \rangle)$ ,  $q(C, \langle C, H \rangle)$
  3. **Solution:**
    - Assuming P1 starts at  $A$ , then P2 goes at  $C$ , then P1 goes at  $H$ , that means after 12 transitions (**even number**), P1 is the last turn-taker, therefore, P1 gets the reward of +4.
    - **Backpropagation:**
      - $N(C, \langle C, H \rangle) = 1$ ,  $X(C, \langle C, H \rangle) = 4$  so 4/1
      - $N(A, \langle A, C \rangle) = 4$ ,  $X(A, \langle A, C \rangle) = 2 + 4 = 6$  so 6/4
      - $q(A, \langle A, B \rangle) = 2/5 = 0.4$
      - $q(A, \langle A, C \rangle) = 6/4 = 1.5$
      - $q(C, \langle C, H \rangle) = 4/1 = 4$